

Signal Representations

1 Motivation

Fourier representations, including the Fourier series and the DFT, are a pillar of signal processing, but they have a fundamental limitation: Fourier coefficients capture global fluctuations of a signal, and hence do not provide information about events that are localized in time or space. In this chapter we study signal representations that address this issue. Section 2 describes the short-time Fourier transform, designed to capture frequency information that is localized in time. Section 3 introduces wavelet transforms, which decompose signals into components with different resolutions. As explained, in Section 4, these transformations can be used to generate sparse representations of audio and image signals, which are useful for tasks like denoising.

2 Time-frequency representations

2.1 Signal segmentation using windows

The Fourier series and the DFT reveal the periodic components of signals. However, they do not provide any information about the time structure of these periodic components. Consider the audio signal in Figure 1. Like most speech and music signals, the signal consists of periodic oscillations that *change over time*. An effective way of capturing such structure is to compute the Fourier coefficients of a signal after segmenting it into intervals. Segmentation can be achieved through multiplication with a window function, which is nonzero only over a specific interval. The simplest possible choice is the rectangular window.

Definition 2.1 (Rectangular window). *The rectangular window $\vec{\pi} \in \mathbb{C}^N$ with width w is defined as*

$$\vec{\pi}[j] := \begin{cases} 1 & \text{if } |j| \leq w, \\ 0 & \text{otherwise.} \end{cases} \quad (1)$$

The rectangular window does not distort the signal at all within the segment over which it is nonzero. However, we are interested in obtaining a local frequency representation, so we need to evaluate the distortion in the frequency domain. As established in the following theorem, multiplying a signal by a window is equivalent to convolving their DFTs.

Theorem 2.2 (Multiplication in time is convolution in frequency). *Let $y := x_1 \circ x_2$ for $x_1, x_2 \in \mathbb{C}^N$. Then the DFT of y equals*

$$\hat{y} = \frac{1}{N} \hat{x}_1 * \hat{x}_2, \quad (2)$$

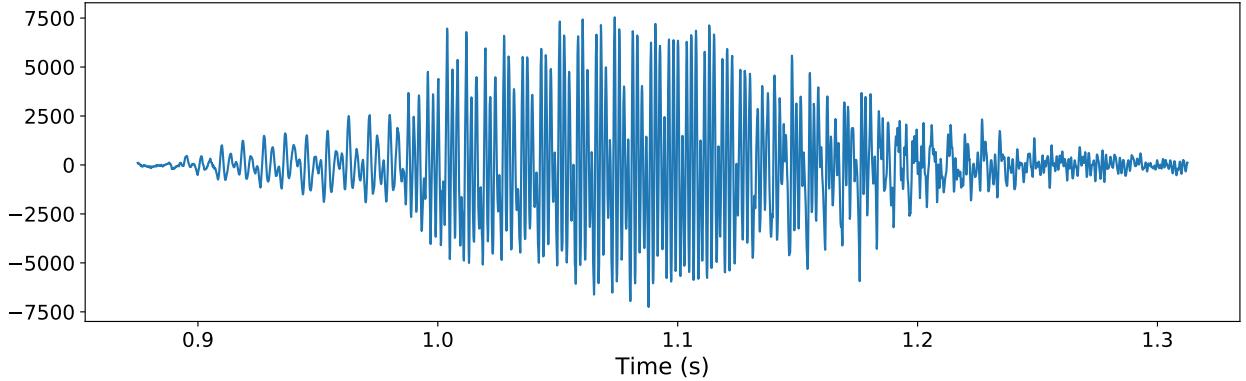


Figure 1: Audio signal of a person saying the word *no* in Hebrew.

\hat{x}_1 and \hat{x}_2 are the DFTs of x_1 and x_2 respectively.

Proof. We have

$$\hat{y}[k] := \sum_{j=1}^N x_1(j)x_2(j) \exp\left(-\frac{i2\pi kj}{N}\right) \quad (3)$$

$$= \sum_{j=1}^N \frac{1}{N} \sum_{l=1}^n \hat{x}_1(l) \exp\left(\frac{i2\pi lj}{N}\right) x_2(j) \exp\left(-\frac{i2\pi kj}{N}\right) \quad (4)$$

$$= \frac{1}{N} \sum_{l=1}^N \hat{x}_1(l) \sum_{j=1}^N x_2(j) \exp\left(-\frac{i2\pi(k-l)j}{N}\right) \quad (5)$$

$$= \frac{1}{N} \sum_{l=1}^N \hat{x}_1(l) \hat{x}_2^{\downarrow l}[k]. \quad (6)$$

□

The following lemma derives the DFT of the rectangular window, which is called a discretized sinc in the signal-processing literature. To simplify the exposition we index the vector and its DFT from $-N/2 + 1$ to $N/2$, assuming that N is even.

Lemma 2.3. *The DFT coefficients of the rectangular window $\vec{\pi} \in \mathbb{C}^N$ with width $2w$ from Definition 2.1*

$$\vec{\pi}[j] := \begin{cases} 1 & \text{if } |j| \leq w, \\ 0 & \text{otherwise,} \end{cases} \quad (7)$$

equal

$$\hat{\pi}[k] = \begin{cases} 2w + 1 & \text{for } k = 0 \\ \frac{\sin\left(\frac{2\pi k(w+1/2)}{N}\right)}{\sin\left(\frac{\pi k}{N}\right)}, & \text{for } k \in \{-N/2 + 1, -N/2 + 2, \dots, N/2\} / 0, \end{cases} \quad (8)$$

Proof. We have

$$\hat{\pi}(0) = \sum_{j=-N/2+1}^{N/2} \vec{\pi}[j] \quad (9)$$

$$= \sum_{j=-w}^w 1 \quad (10)$$

$$= 2w + 1. \quad (11)$$

For $k \neq 0$,

$$\hat{\pi}(k) = \sum_{j=-N/2+1}^{N/2} x[j] \exp\left(-\frac{i2\pi kj}{N}\right) \quad (12)$$

$$= \sum_{j=-w}^w \exp\left(-\frac{i2\pi k}{N}\right)^j \quad (13)$$

$$= \frac{\exp\left(\frac{i2\pi kw}{N}\right) - \exp\left(-\frac{i2\pi k(w+1)}{N}\right)}{1 - \exp\left(-\frac{i2\pi k}{N}\right)} \quad (14)$$

$$= \frac{\exp\left(-\frac{i2\pi k}{2N}\right) 2i \sin\left(\frac{2\pi k(w+1/2)}{N}\right)}{\exp\left(-\frac{i2\pi k}{2N}\right) 2i \sin\left(\frac{\pi k}{N}\right)} \quad (15)$$

$$= \frac{\sin\left(\frac{2\pi k(w+1/2)}{N}\right)}{\sin\left(\frac{\pi k}{N}\right)}. \quad (16)$$

□

Unfortunately, the sinc function has side lobes and does not decay rapidly. As a result, convolving the DFT of a signal by the DFT of a rectangular window produces significant distortion in the frequency domain, as shown in Figure 2. The reason is the discontinuity in the rectangular window. It introduces an artificial discontinuity in the windowed function, which gives rise to spurious high-frequency components. The solution is to use a tapered window, which decreases smoothly to zero at the borders. A popular choice is the Hann window. As we can see in Figure 2, this window produces much less distortion in the frequency domain.

Definition 2.4. *The Hann window $h \in \mathbb{R}^N$ of width $2w$ equals*

$$h[j] := \begin{cases} \frac{1}{2} (1 + \cos(\frac{\pi j}{w})) & \text{if } |j| \leq w, \\ 0 & \text{otherwise.} \end{cases} \quad (17)$$

Multiplying by a Hann window (or other similar options) and then computing the DFT is a popular technique for analyzing the frequency content of signals locally. The time resolution of the analysis is governed by the width of the window. Reducing the window therefore increases the resolution in time. However, there is a price to pay. The frequency resolution of the analysis is governed by the width of the window in the frequency domain. The following theorem shows that compressing a signal in time, dilates it in frequency, and vice versa.

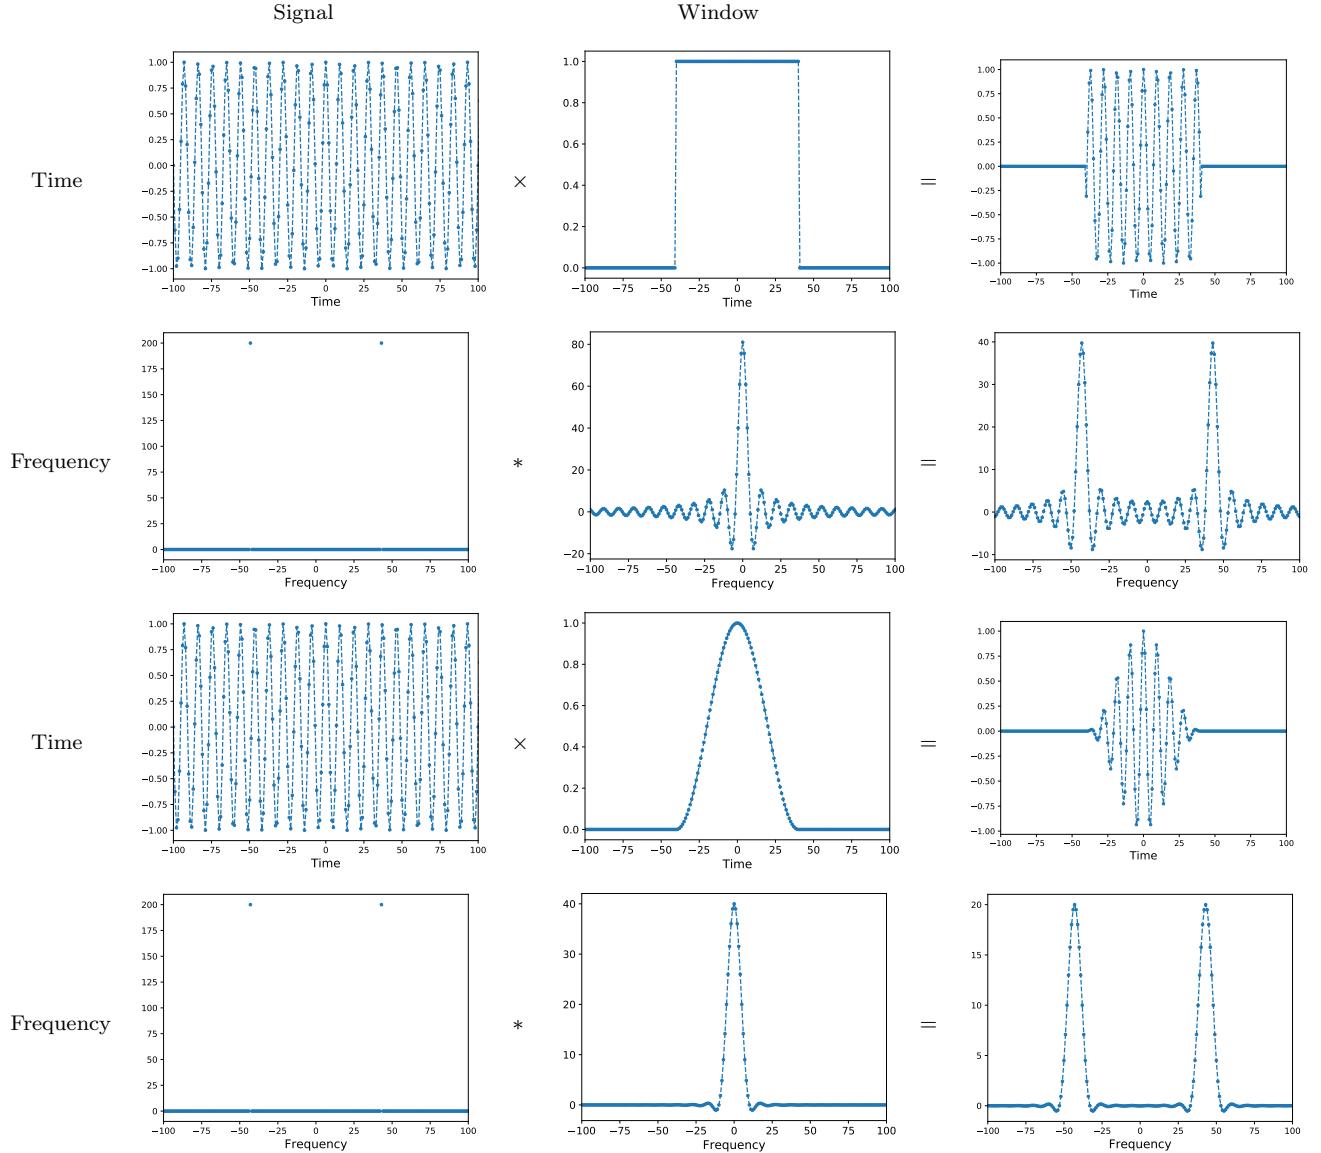


Figure 2: The right column shows the results of multiplying a sinusoidal signal (left column) by a rectangular or a Hann window (central column) in the time domain.

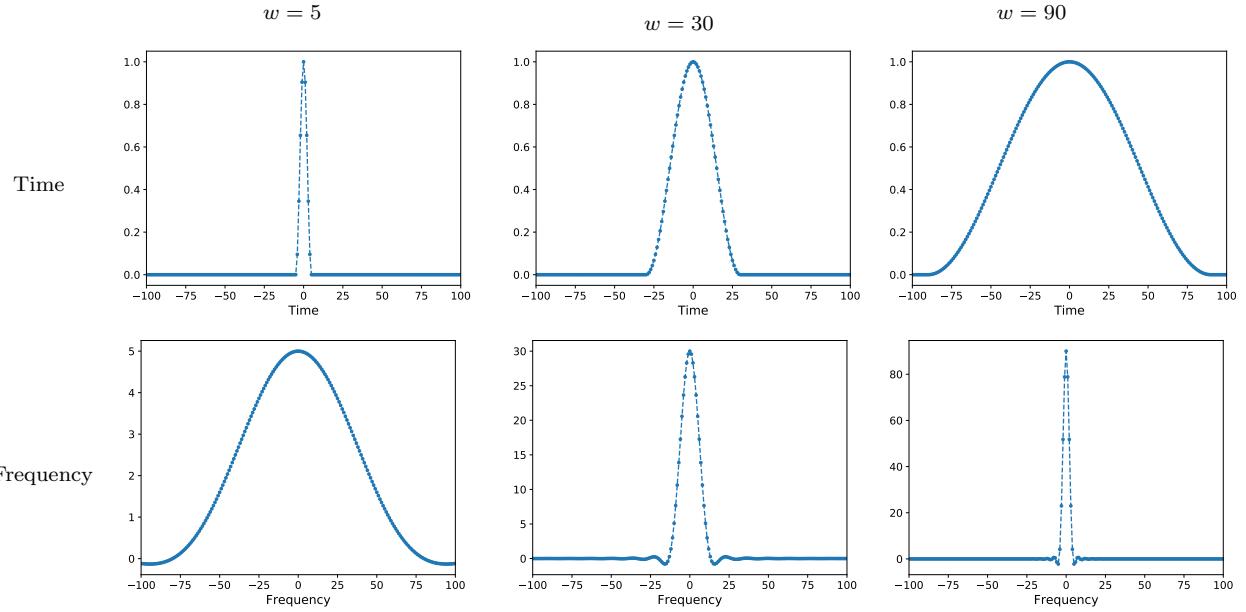


Figure 3: Dilations and contractions of a Hann window in time and the corresponding effect on the DFT of the window.

Theorem 2.5. Let $x \in \mathcal{L}_2[-T/2, T/2]$, $T > 0$, which is nonzero only in a band of width $2w$ around zero, i.e. $x(t) = 0$ for $t > |w|$ where $w > 0$. The Fourier series coefficients of the signal y that satisfies

$$y(t) = x(\alpha t), \quad \text{for all } t \in [-T/2, T/2], \quad (18)$$

equal

$$\hat{y}[k] = \frac{1}{\alpha} \langle x, \phi_{k/\alpha} \rangle \quad (19)$$

for any positive real number α such that $w/\alpha < T/2$.

Proof. By definition of the Fourier series coefficients

$$\hat{y}[k] = \int_{t=-T/2}^{T/2} y(t) \exp\left(-\frac{i2\pi kt}{T}\right) dt \quad (20)$$

$$= \int_{t=-w/\alpha}^{w/\alpha} x(\alpha t) \exp\left(-\frac{i2\pi kt}{T}\right) dt \quad (21)$$

$$= \frac{1}{\alpha} \int_{\tau=-w}^w x(\tau) \exp\left(-\frac{i2\pi k\tau}{\alpha T}\right) d\tau \quad (22)$$

$$= \frac{1}{\alpha} \int_{\tau=-T/2}^{T/2} x(\tau) \exp\left(-\frac{i2\pi k\tau}{\alpha T}\right) d\tau. \quad (23)$$

□

There is consequently a fundamental trade-off between the resolution in time and frequency that we can achieve simultaneously. If we make a window narrower to increase the time resolution, its DFT will become wider. As a result, the convolution in the frequency domain will blur the frequency content of the signal to a greater extent, decreasing the frequency resolution of the analysis. This fundamental tradeoff is known as the *uncertainty principle*. We refer the interested reader to Section 2.3.2 of Ref. [1] for more details.

2.2 Short-time Fourier transform

As discussed in the previous section, frequency representations such as the Fourier series and the DFT provide global information about the fluctuations of a signal, but they do not capture *local* information. The short-time Fourier transform (STFT) is designed to describe localized fluctuations. The STFT coefficients are equal to the DFT of time segments of the signal, extracted through multiplication with a window. Here we focus on the discrete STFT; it is also possible to define continuous versions.

Definition 2.6 (Short-time Fourier transform). *The short-time Fourier transform of a vector $x \in \mathbb{C}^N$ is given by*

$$\text{STFT}_{[\ell]}(x)[k, s] := \left\langle x, \xi_k^{\downarrow s(1-\alpha_{\text{ov}})\ell} \right\rangle, \quad 0 \leq k \leq \ell - 1, \quad 0 \leq s \leq \frac{N}{(1-\alpha_{\text{ov}})\ell}, \quad (24)$$

where the basis vectors are defined as

$$\xi_k[j] := \begin{cases} w_{[\ell]}(j) \exp\left(\frac{i2\pi kj}{\ell}\right) & \text{if } 1 \leq j \leq \ell \\ 0 & \text{otherwise.} \end{cases} \quad (25)$$

The overlap between adjacent segments equals $\alpha_{\text{ov}}\ell$. The STFT is parametrized by the choice of segment length ℓ , window function $w_{[\ell]} \in \mathbb{C}^\ell$, and overlap factor α_{ov} .

The STFT coefficients are defined as the inner product between the signal and basis vectors obtained by shifting the window both in time and in frequency. The frequency shifts are achieved through multiplication with complex sinusoids (as you will prove in the homework). Figure 4 shows the different basis vectors for a specific example.

Due to the overlap between the shifted signals in time, the STFT is an *overcomplete* transformation. If the overlap factor $\alpha_{\text{ov}} := 0.5$ then there the dimensionality of the STFT is approximately double (up to border effects) that of the original vector. This is apparent in the following matrix representation of the STFT,

$$\begin{bmatrix} \text{STFT}_{[\ell]}(x)[0, 0] \\ \dots \\ \text{STFT}_{[\ell]}(x)[\ell - 1, 0] \\ \text{STFT}_{[\ell]}(x)[0, 1] \\ \dots \\ \text{STFT}_{[\ell]}(x)[\ell - 1, 1] \\ \text{STFT}_{[\ell]}(x)[0, 2] \\ \dots \\ \text{STFT}_{[\ell]}(x)[\ell - 1, 2] \\ \dots \end{bmatrix} := \begin{bmatrix} F_{[\ell]} & 0 & 0 & 0 & \dots \\ 0 & F_{[\ell]} & 0 & 0 & \dots \\ 0 & 0 & F_{[\ell]} & 0 & \dots \\ 0 & 0 & 0 & F_{[\ell]} & \dots \\ 0 & 0 & 0 & 0 & \dots \\ \dots & \dots & \dots & \dots & \dots \end{bmatrix} \begin{bmatrix} \text{diag}(w_{[\ell]}) & 0 & 0 & \dots \\ 0 & \text{diag}(w_{[\ell]}) & 0 & \dots \\ 0 & 0 & \text{diag}(w_{[\ell]}) & \dots \\ 0 & 0 & 0 & \dots \\ \dots & \dots & \dots & \dots \end{bmatrix} x,$$

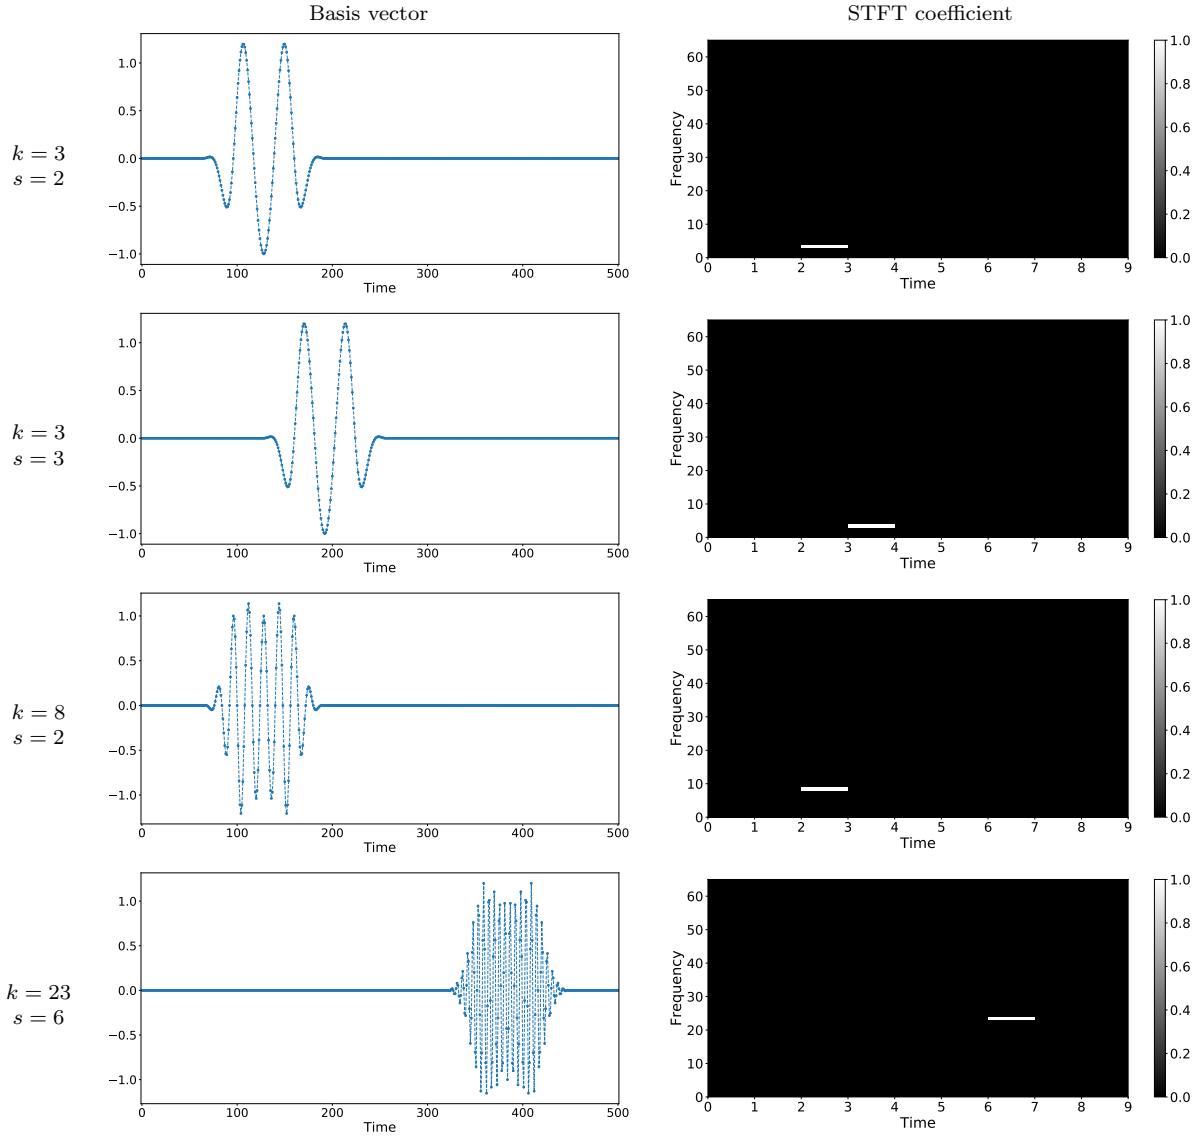


Figure 4: The left columns shows different basis functions of the STFT for a signal of length $N := 500$ when $\ell := 128$ and $\alpha_{\text{ov}} := 0.5$ and the window is a Hann window. The right column shows the corresponding STFT coefficient. Only the positive frequency axis is shown.

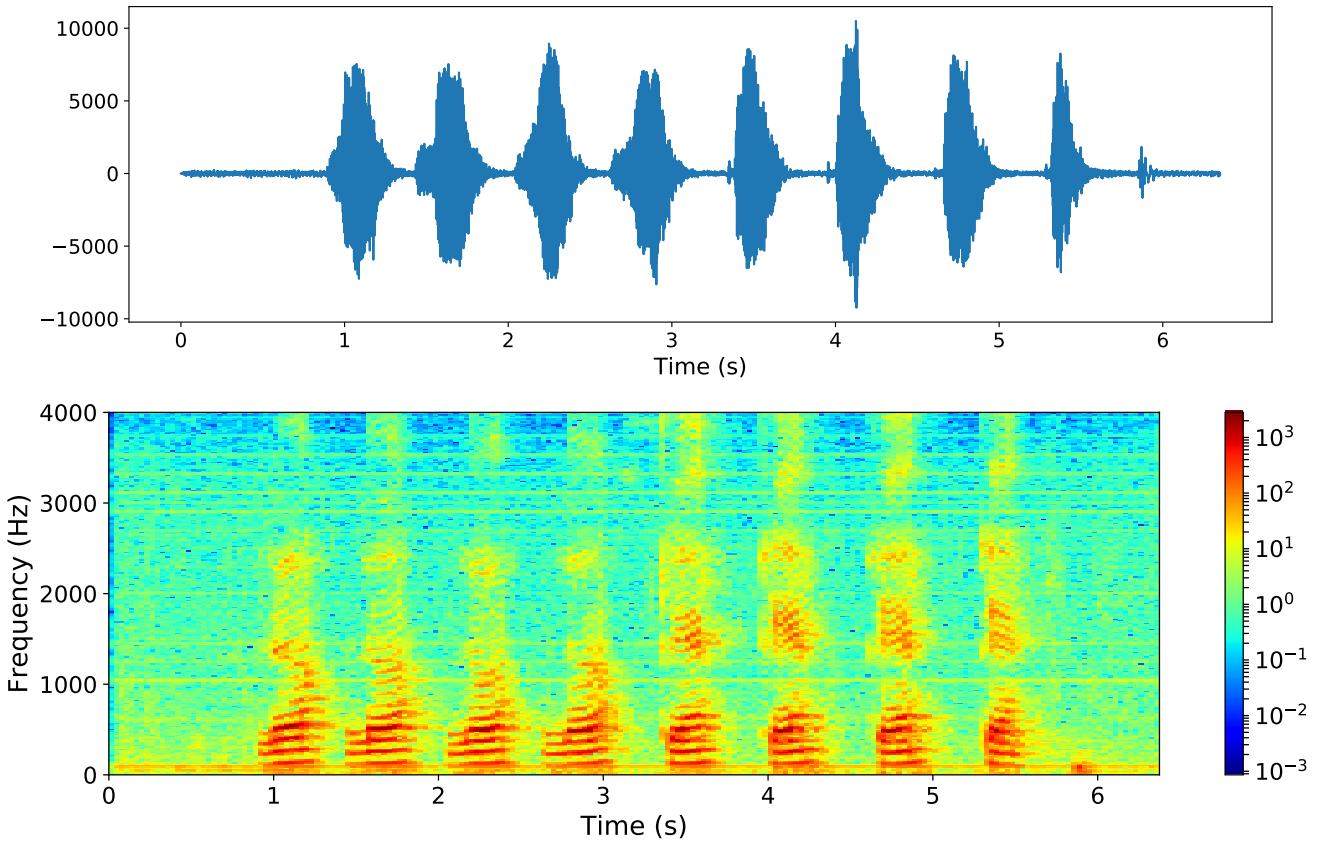


Figure 5: The top image shows the time representation of an audio signal where someone is saying *no* four times and then *yes* four times in Hebrew. The heatmap below shows the magnitude of the STFT coefficients of the signal computed using Hann windows of length 62.5 ms. Only the positive frequency axis is shown.

where $F_{[\ell]}$ is an $\ell \times \ell$ DFT matrix, 0 represents a $\ell/2 \times \ell/2$ matrix full of zeros, and $\text{diag}(w_{[\ell]})$ is a diagonal matrix that has the window function as its diagonal. Overcomplete representations are known as *frames* in signal processing.

The STFT can be computed very efficiently using the FFT algorithm to obtain the DFT of the signal segments. The complexity of multiplying a segment by a window and applying the FFT is $O(\ell \log \ell)$. The number of segments is approximately $N/(1 - \alpha_{\text{ov}})\ell$ (ignoring the borders). The complexity of computing the STFT of a vector of length N is therefore $O(N \log \ell)$ (the overlap factor is usually equal to 0.5 or a similar fraction). To recover a signal from its STFT coefficients we can compute the inverse DFT of each segment, again applying the FFT algorithm, and then combine the scaled segments. The complexity is also $O(N \log \ell)$.

Figure 5 shows the magnitude of the STFT coefficients of a real-world speech signal. This representation, known as the *spectrogram*, is a fundamental tool in audio analysis. It is used to visualize how the signal energy in each frequency band evolves over time.

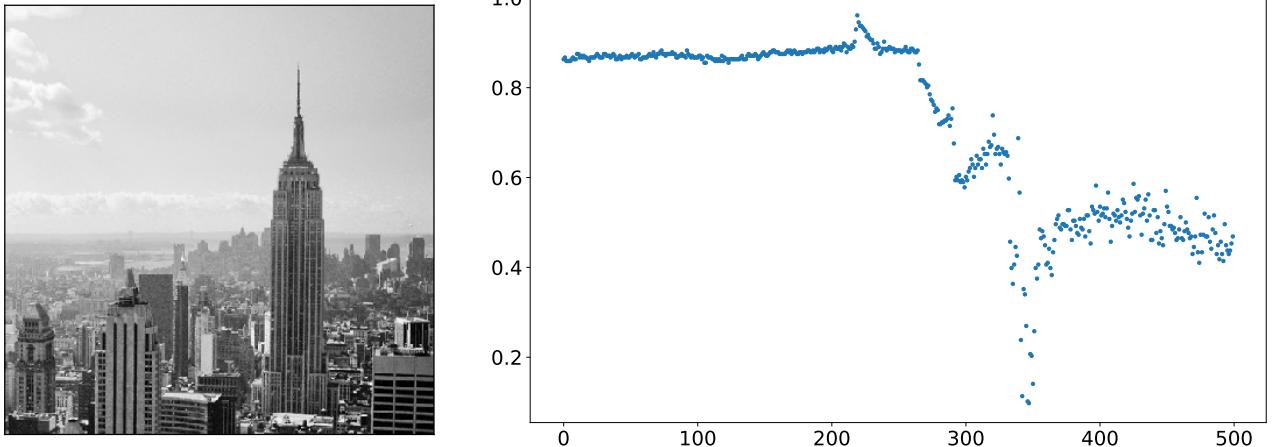


Figure 6: An image of New York City and the values of a subset of pixels situated on a vertical line (if the image is interpreted as a 500×500 image the pixels correspond to column 135).

3 Multiresolution analysis

3.1 Wavelets

The scale at which information is encoded in signals is usually not uniform. For example, the top half of the image in Figure 6 is very smooth, whereas the bottom half has a lot of high-resolution details. The aim of multiresolution analysis is to reveal this structure, by decomposing signals into components corresponding to different scales. To make this precise, we define a decomposition of the ambient signal space into subspaces, each of which encodes information at a certain resolution. This is a discrete version of the framework introduced by Mallat and Meyer (see Chapter 7 in [1]).

Definition 3.1 (Multiresolution decomposition). *Let $N := 2^K$ for some integer K . A multiresolution decomposition of \mathbb{R}^N is a sequence of nested subspaces $\mathcal{V}_K \subset \mathcal{V}_{K-1} \subset \dots \subset \mathcal{V}_0$. The subspaces must satisfy the following properties:*

- $\mathcal{V}_0 = \mathbb{R}^N$, the approximation at scale $2^0 = 1$ is perfect.
- \mathcal{V}_k is invariant to translations of scale 2^{k+1} for $0 \leq k \leq K$. If $x \in \mathcal{V}_k$ then

$$x^{\downarrow l 2^{k+1}} \in \mathcal{V}_k \quad \text{for all } l \in \mathbb{Z}. \quad (26)$$

- Dilating vectors in \mathcal{V}_j by 2 yields vectors in \mathcal{V}_{j+1} . Let $x \in \mathcal{V}_j$ be nonzero only between 1 and $N/2$, the dilated vector $x_{\leftrightarrow 2}$ belongs to \mathcal{V}_{j+1} (see Def. 3.2 below).

Definition 3.2 (Discrete dilation). *Let vector $x \in \mathbb{R}^N$, with entries indexed between 0 and $N - 1$, satisfy $x[j] = 0$ for all $j \geq N/M$, where M is a fixed positive integer. We define the dilation of x by a factor of M as*

$$x_{\leftrightarrow M}[j] = x \left[\left\lceil \frac{j}{M} \right\rceil \right]. \quad (27)$$

Intuitively, subspace \mathcal{V}_j contains information encoded at a scale 2^j . If we dilate a signal in \mathcal{V}_j by a factor of 2, thereby doubling its scale and halving its resolution, then it is guaranteed to belong to \mathcal{V}_{j+1} . By projecting a signal onto \mathcal{V}_j we obtain an approximation at the corresponding resolution.

Mallat and Meyer suggested the following strategy to build multiresolution decompositions:

- Set the coarsest subspace to be spanned by a low-frequency vector φ , called a scaling vector or father wavelet:

$$\mathcal{V}_K := \text{span}(\varphi). \quad (28)$$

- Decompose the finer subspaces into the direct sum

$$\mathcal{V}_k := \mathcal{W}_k \oplus \mathcal{V}_{k+1}, \quad 0 \leq k \leq K-1, \quad (29)$$

where \mathcal{W}_k captures the finest resolution available at level k . Set \mathcal{W}_k to be spanned by shifts of a vector μ dilated to have the appropriate resolution:

$$\mathcal{V}_k := \mathcal{W}_k \oplus \mathcal{V}_{k+1}, \quad 0 \leq k \leq K-1, \quad (30)$$

$$\mathcal{W}_k := \bigoplus_{m=0}^{\frac{N-1}{2^{k+1}}} \text{span}\left(\mu_{\leftrightarrow 2^k}^{\downarrow m 2^{k+1}}\right). \quad (31)$$

The vector μ is called a mother wavelet.

Each subspace is spanned by scaled and shifted copies of the mother wavelet, except for the coarsest resolution, which is spanned by the father wavelet (alternatively, it could be spanned by shifted copies of the father wavelet). Designing valid multiresolution decompositions requires selecting a specific father and mother wavelet pair, such that the corresponding subspaces satisfy the requirements. The simplest example is the Haar wavelet.

Definition 3.3 (Haar wavelet basis). *The Haar father wavelet $\varphi \in \mathbb{R}^N$ is a constant vector, such that*

$$\varphi[j] := \frac{1}{\sqrt{N}}, \quad 1 \leq j \leq N. \quad (32)$$

The mother wavelet $\mu \in \mathbb{R}^N$ satisfies

$$\mu[j] := \begin{cases} -\frac{1}{\sqrt{2}}, & j = 1, \\ \frac{1}{\sqrt{2}}, & j = 2, \\ 0, & j > 2. \end{cases} \quad (33)$$

One can easily check that the Haar wavelet provides a valid multiresolution decomposition. The basis vectors are all orthogonal and unit norm, so they form an orthonormal basis of \mathbb{R}^N . Figure 7 shows the basis vectors for $N = 8$. In Figure 8 the decomposition is applied to obtain a multiscale decomposition of the 1D signal corresponding to a vertical line of the image in Figure 6. Approximations obtained using the Haar wavelet are piecewise constant, which may not be desirable. It is

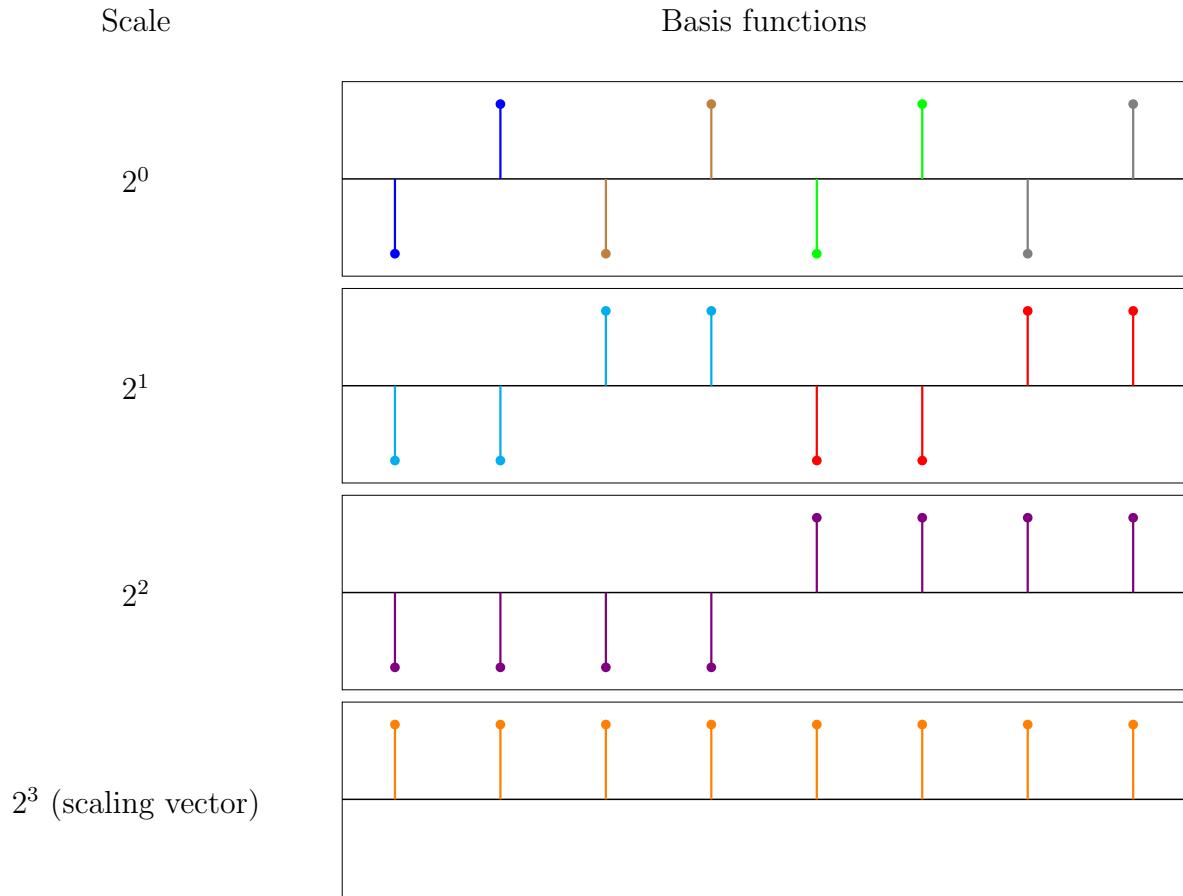


Figure 7: Basis vectors in the Haar wavelet basis for \mathbb{R}^8 . Each row shows the vectors corresponding to the corresponding scale in different colors.

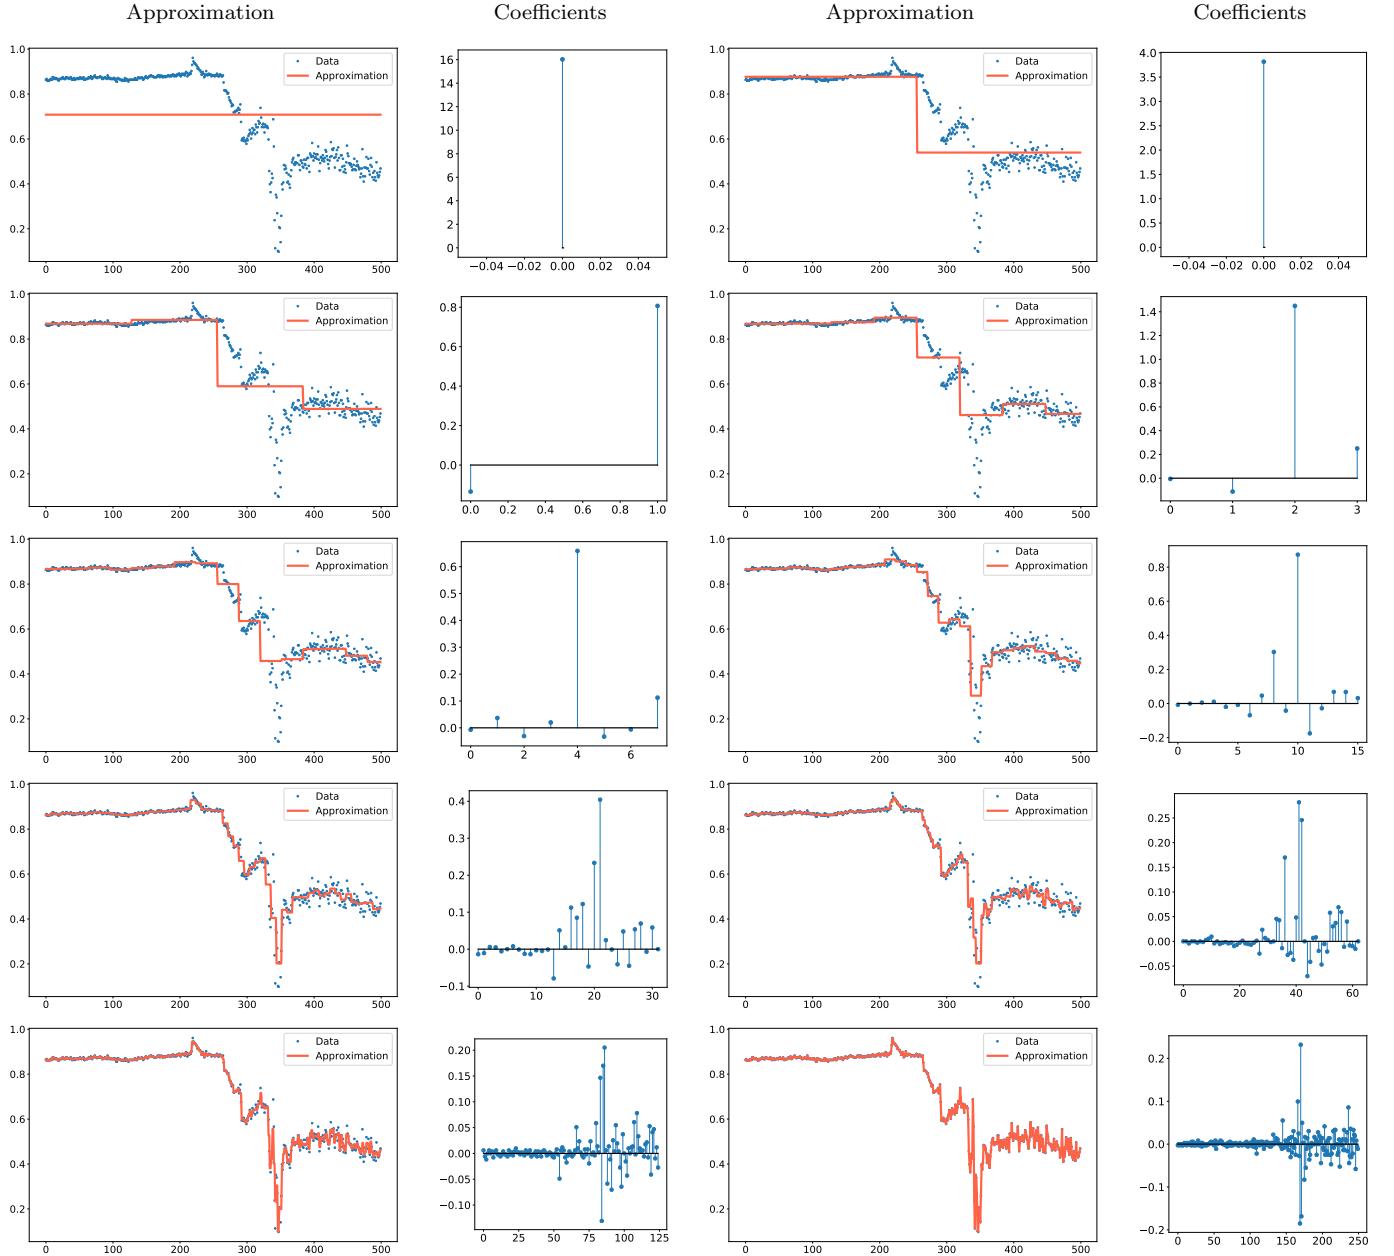


Figure 8: Multiresolution decomposition of the 1D signal corresponding to a vertical line of the image in Figure 6 using the 1D Haar wavelet basis. The corresponding wavelet coefficients are plotted on the right for each scale.

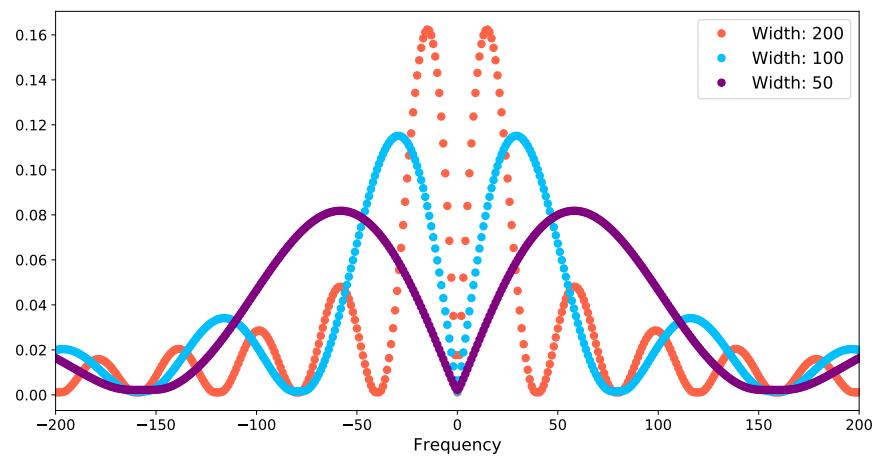


Figure 9: DFT of Haar wavelets with different widths.

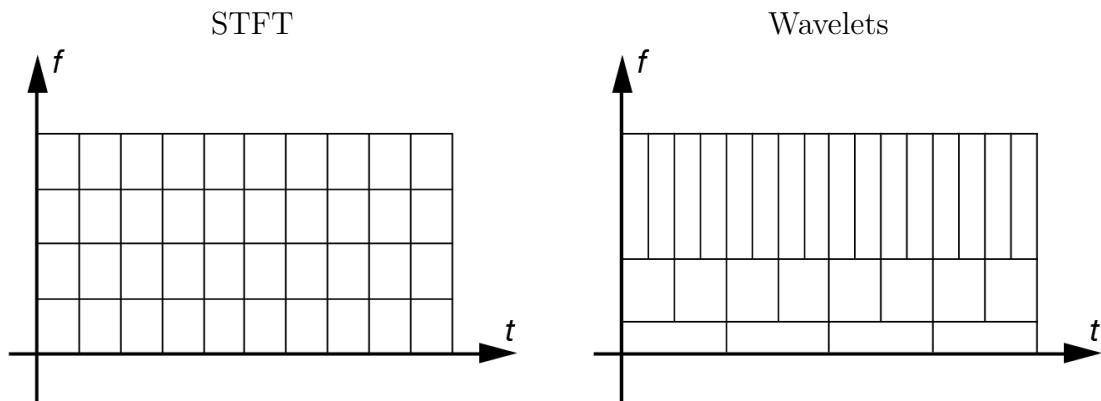


Figure 10: Diagram of the time-frequency support of STFT (left) and wavelet basis vectors (right).

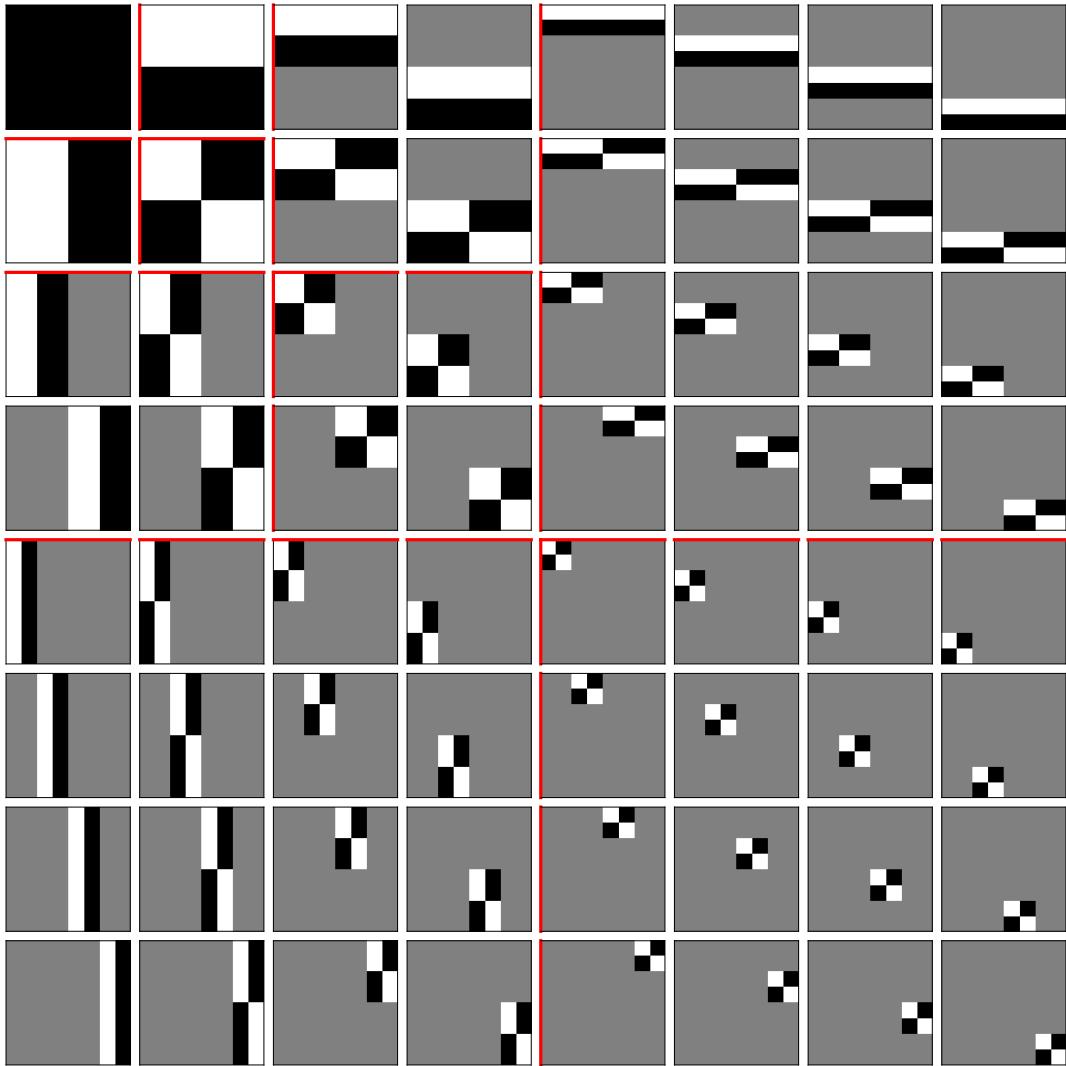


Figure 11: Basis vectors of the 2D Haar wavelet decomposition.

also possible to build multiresolution decompositions with smoother wavelets. Examples include Meyer wavelets, Daubechies wavelets, coiflets, symmlets, etc. We refer the interested reader to Chapter 7 in [1] for a detailed description.

In contrast to the STFT, multiresolution analysis provides a time-frequency decomposition of signals where different coefficients have different resolutions. Figure 9 shows the DFT of Haar wavelets corresponding to different scales. As established in Theorem 2.5, when the wavelet contracts (decreasing the time resolution), its DFT dilates (increasing the frequency resolution). Note that the Haar wavelet is not very localized in frequency due to its discontinuity, just like the rectangular window (see Lemma 2.3). Figure 10 provides a cartoon illustration of the time-frequency supports of STFT and wavelet basis vectors. The STFT tiles the time-frequency plane regularly, whereas wavelet vectors have very different time and frequency resolutions.

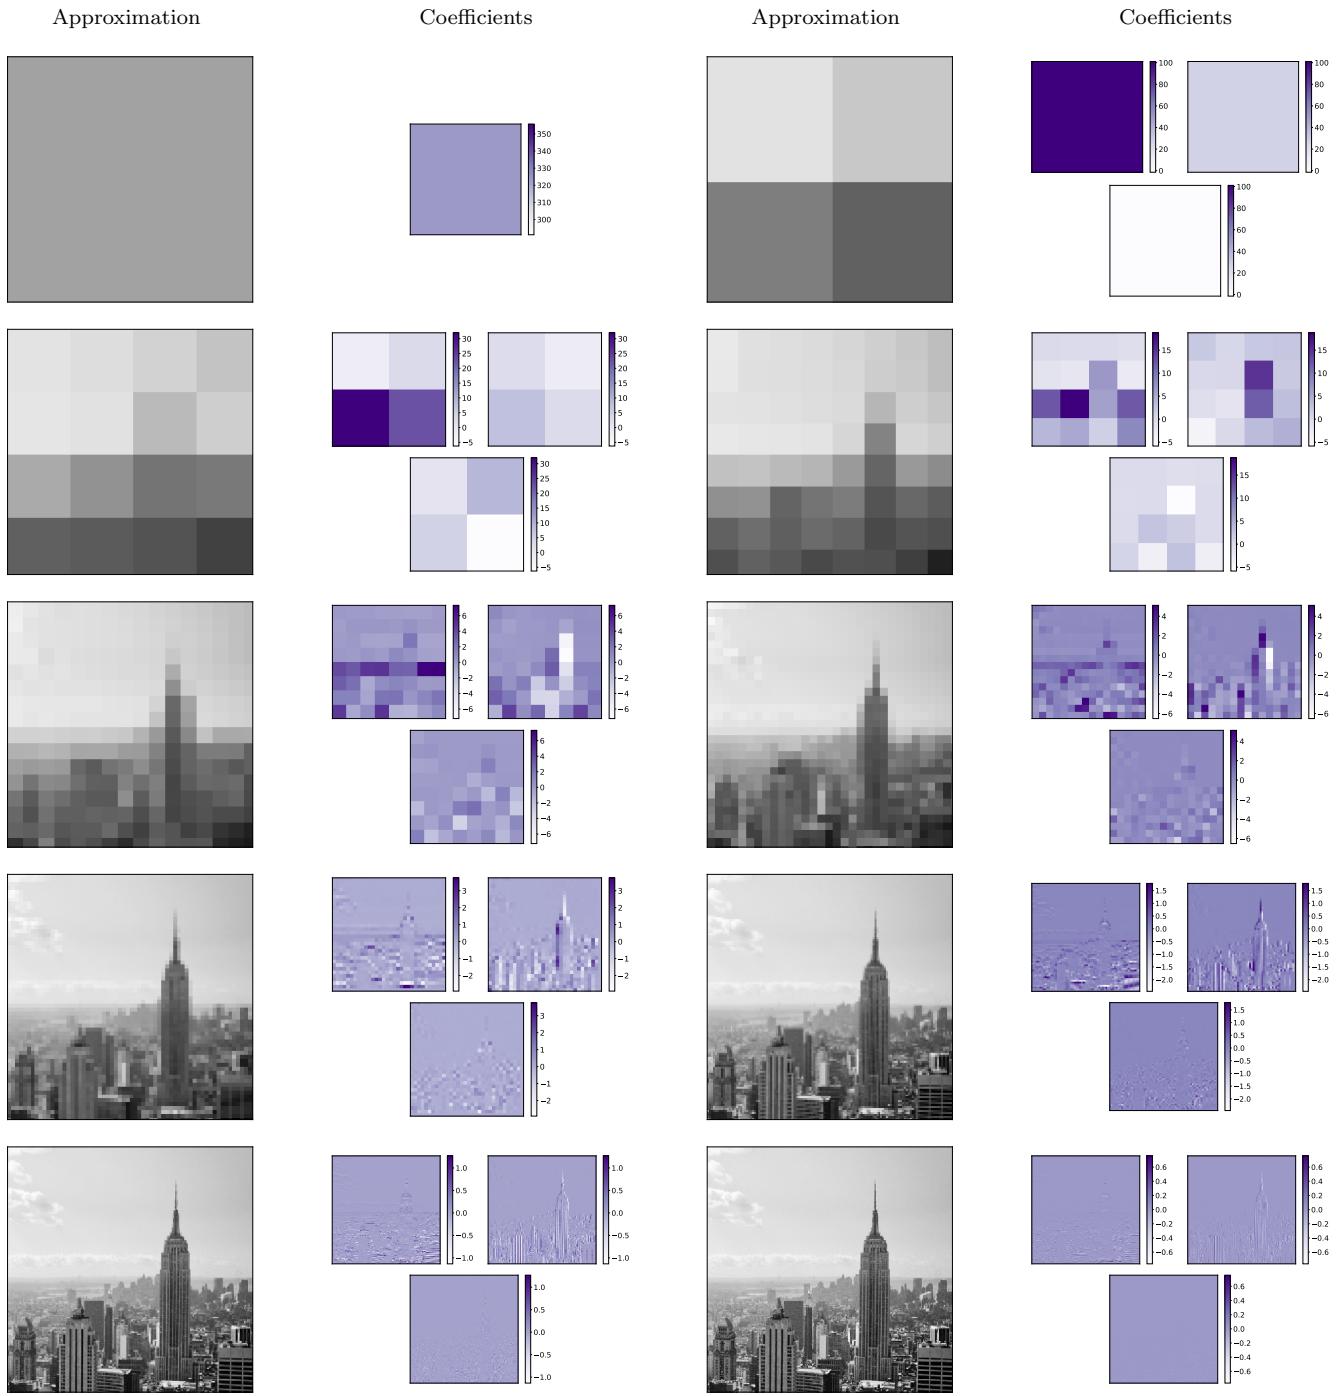


Figure 12: Multiresolution decomposition of the image in Figure 6 using the 2D Haar wavelet basis. The corresponding wavelet coefficients are plotted on the right for each scale.

3.2 Multidimensional wavelets

Multiresolution decompositions are an important tool in image processing, where they are applied to 2D and 3D signals. A simple way to design multidimensional decompositions is to leverage 1D multiresolution decompositions. The basis vectors of each subspace are constructed by taking outer products of the corresponding one-dimensional wavelets. Let $v_{[m,s]}^{1D}$ denote a basis vector at scale m shifted by s (depending on the resolution it can be a father or mother wavelet). To build a 2D basis vector at scale (m_1, m_2) and shift (s_1, s_2) we set

$$v_{[s_1, s_2, m_1, m_2]}^{2D} := v_{[s_1, m_1]}^{1D} (v_{[s_2, m_2]}^{1D})^T, \quad (34)$$

where v^{1D} can refer to 1D father or mother wavelets.

Figure 11 shows the resulting 2D basis vectors for the Haar wavelet basis. Two-dimensional wavelet transforms are a popular tool to perform multiresolution decompositions of images. Figure 12 shows a multiresolution decomposition of the image in Figure 6 using the 2D Haar wavelet basis. One can use any 1D wavelets to build a 2D multiresolution decomposition. However, better performance can be achieved by designing non-separable basis vectors. Examples of non-separable 2D wavelets include the steerable pyramid, curvelets, and bandlets. See [here](#) and Section 9.3 in [1] for more details.

4 Denoising via thresholding

4.1 Hard thresholding

The STFT and wavelet transforms often yield sparse representations of signals, where many coefficients are close to zero. In the case of the STFT, this occurs when only a few spectral components are active at a particular time, which is typical of speech or music signals (see Figure 5). In the case of wavelets, sparsity results from the fact that large regions of natural images (and many other signals) are smooth and mostly contain coarse-scale features, whereas most of the fine-scale features are confined to edges or regions with high-frequency textures (see Figure 12).

In contrast, noisy perturbations usually have dense coefficients in any fixed frame or basis. A linear transformation $A\tilde{z}$ of a Gaussian vector with mean μ and covariance matrix Σ is still Gaussian with mean $A\mu$ and covariance matrix $A\Sigma A^*$. If A is an orthogonal matrix and the noise is iid Gaussian with mean zero and variance σ^2 , then $A\tilde{z}$ is also iid Gaussian with mean zero and variance σ^2 (the covariance matrix equals $A\sigma^2 I A^* = N\sigma^2 I$). This means, for example, that the Haar wavelet coefficients of the noise will not be sparse like the ones of a natural image.

Let us consider the problem of denoising measurements $y \in \mathbb{C}^n$ of a signal $x \in \mathbb{C}^n$ corrupted by additive noise $z \in \mathbb{C}^n$

$$y := x + z. \quad (35)$$

Under the assumptions that (1) Ax is sparse representation where A contains the basis vectors of a basis or a frame, and (2) the entries of Az are small and dense, eliminating small entries in the

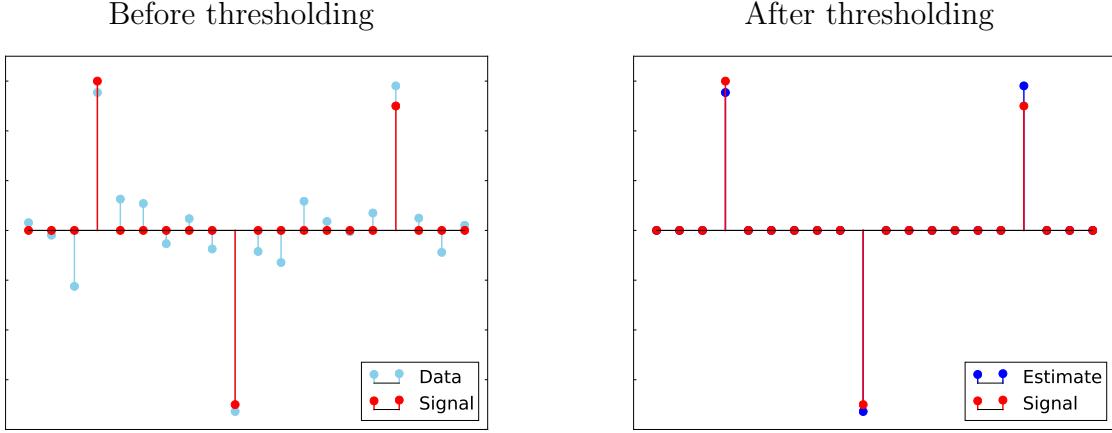


Figure 13: Denoising via hard thresholding.

coefficients

$$Ay = Ax + Az \quad (36)$$

should therefore suppress the noise while preserving the signal. Figure 13 shows a simple example, where the signal itself is sparse (A is just the identity matrix).

Algorithm 4.1 (Denoising via hard thresholding). *Let y follow the model in equation (35) and let A be a full-rank linear transformation that sparsifies the signal x . To denoise we:*

1. *Compute the coefficients Ay .*
2. *Apply the hard-thresholding operator $\mathcal{H}_\eta : \mathbb{C}^n \rightarrow \mathbb{C}^n$ to the coefficients. Let $v \in \mathbb{C}^N$, the operator is defined as*

$$\mathcal{H}_\eta(v)[j] := \begin{cases} v[j] & \text{if } |v[j]| > \eta, \\ 0 & \text{otherwise,} \end{cases} \quad (37)$$

for $1 \leq j \leq N$. The threshold η can be adjusted according to the standard deviation of Az , or by cross validation.

3. *Compute the estimate by inverting the transform, i.e. setting*

$$x_{\text{est}} := L \mathcal{H}_\eta(Ay), \quad (38)$$

where L is a left inverse of A .

Recall that in Wiener filtering the denoising procedure is linear and therefore does not depend on the specific input. Thresholding is nonlinear: different inputs are scaled by different factors (1 or 0). This makes it possible to adapt to the local structure of the signal and usually results in more effective denoising. Figures 14, 15 and 16 show the results of denoising a speech signal applying hard thresholding. Figures 17 and 18 show the results of denoising an image by thresholding its 2D Haar wavelet coefficients. In both cases the threshold is chosen by cross validation on a separate set of signals.

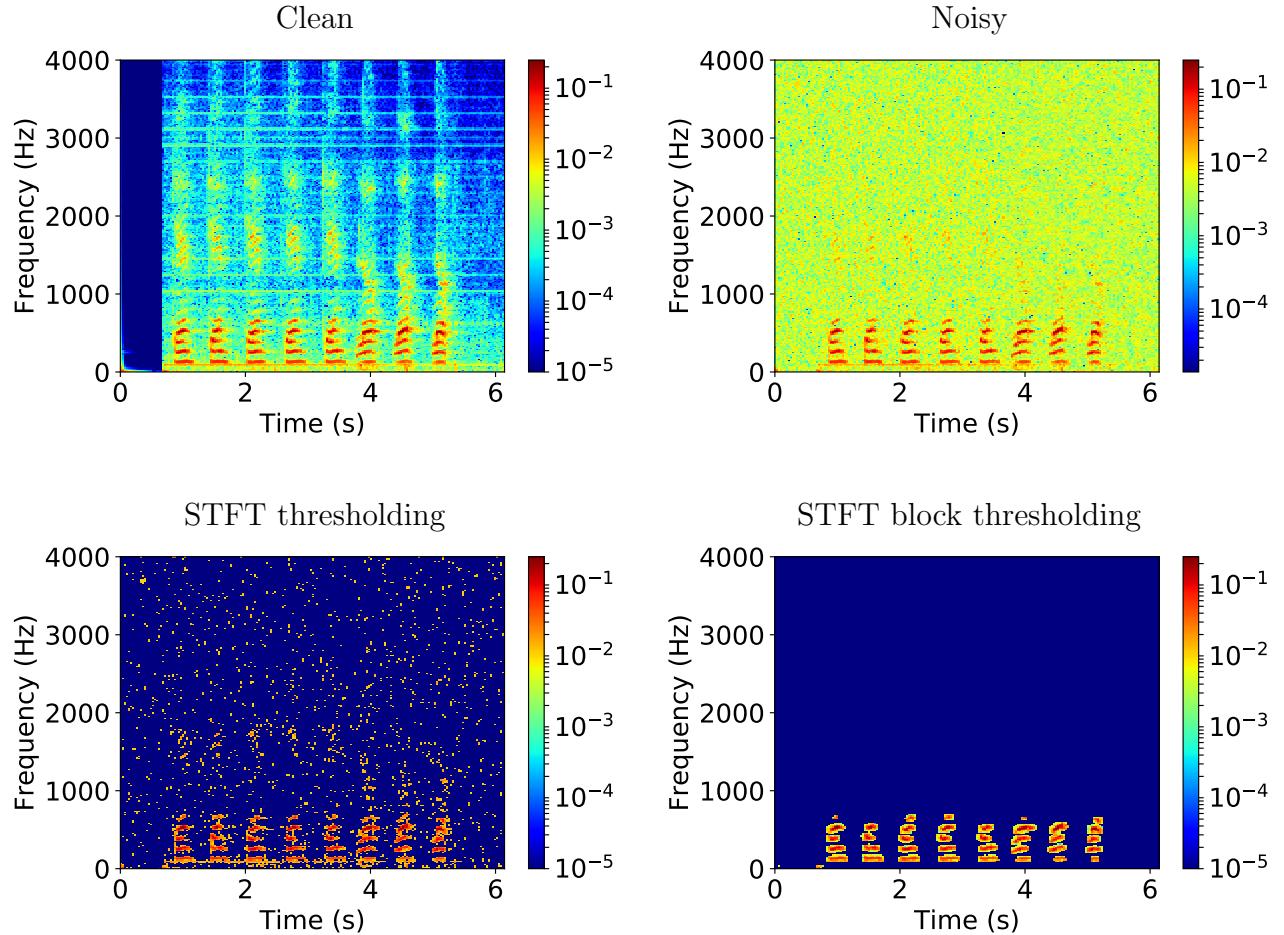


Figure 14: The top row shows the spectrogram of a test audio signal before and after being corrupted by additive iid Gaussian noise with standard deviation $\sigma = 0.1$. The bottom row shows the spectrogram of the signal denoised via STFT thresholding and block thresholding. Only the positive frequency axis is shown.

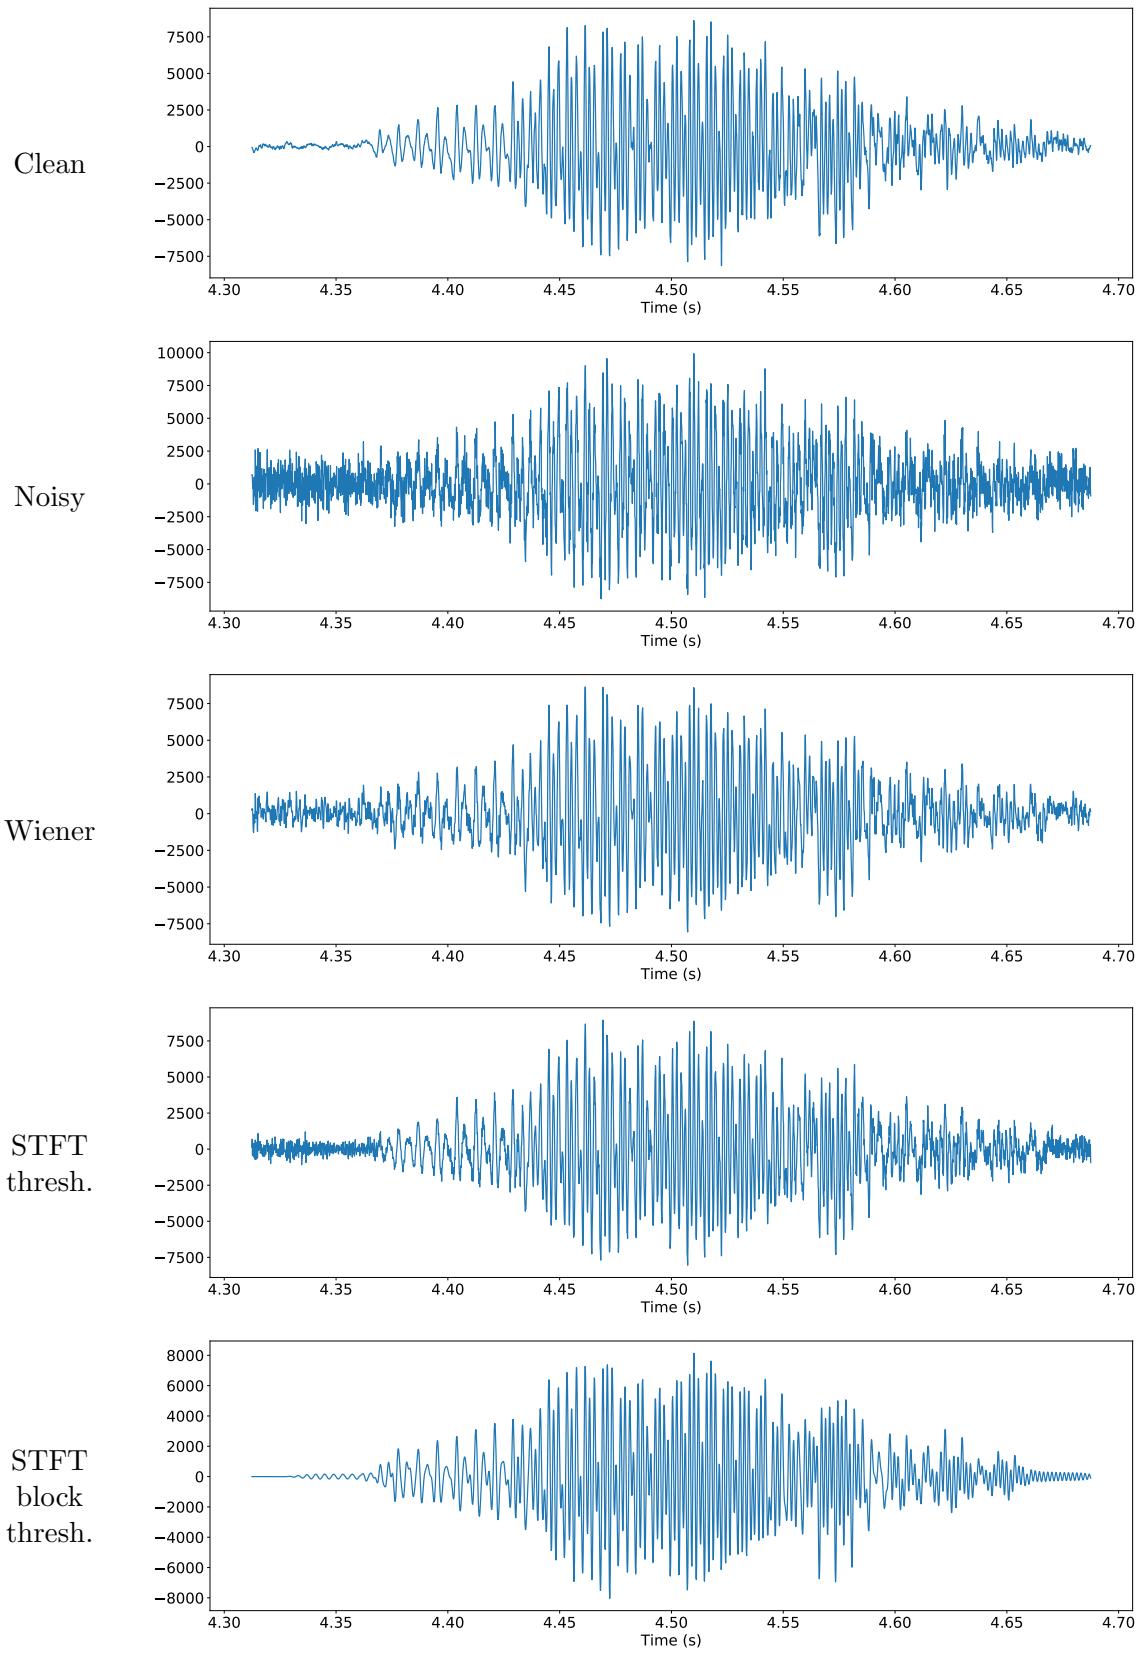


Figure 15: The top two images show the clean and noisy signals in Figure 14 in the time domain. Below, the signal denoised by Wiener filtering is compared to the result of applying STFT thresholding and block thresholding. Click on these links to listen to the audio: [clean signal](#), [noisy signal](#), denoised signal obtained by [Wiener filtering](#), [STFT thresholding](#), [STFT block thresholding](#).

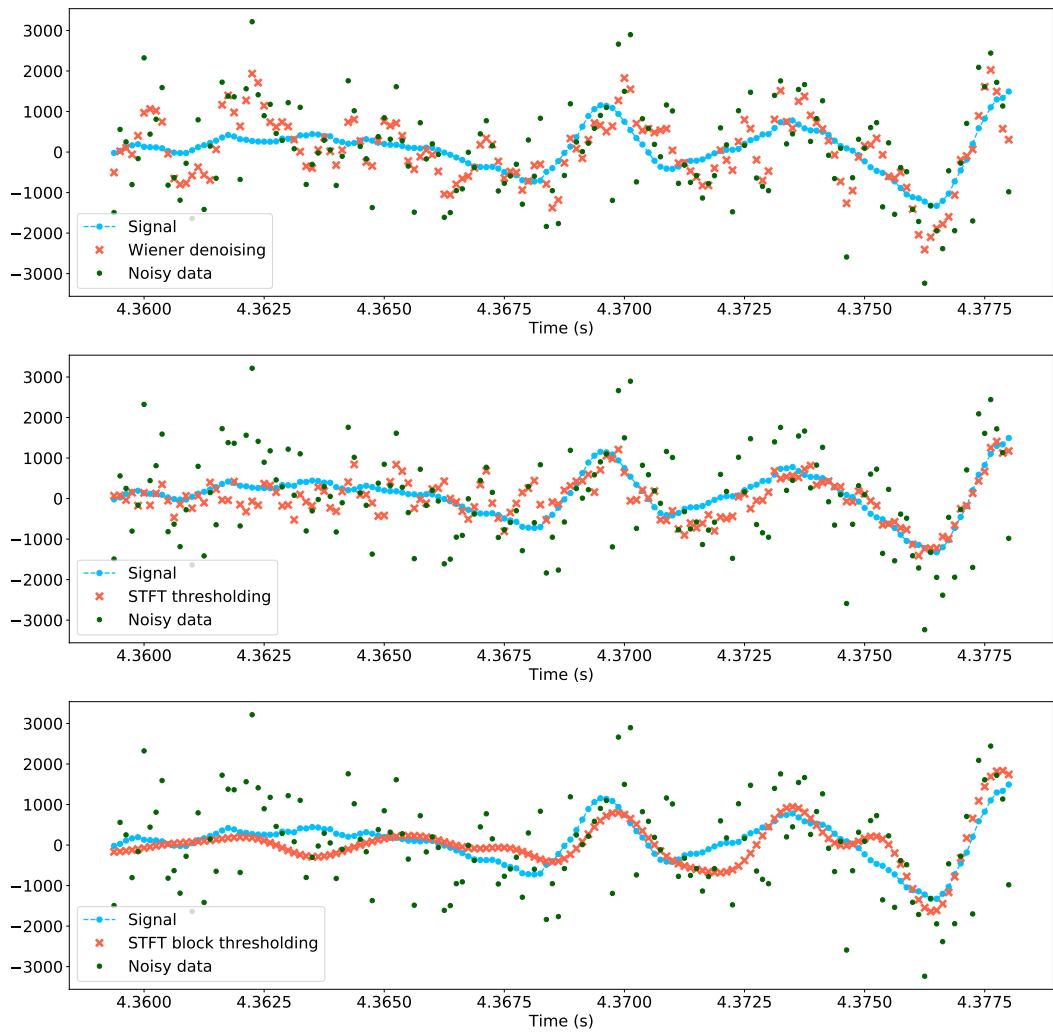


Figure 16: Zoomed-in plots of the signals in Figure 15.

4.2 Block thresholding

When we compute representations that capture localized details of signals, such as the wavelet transform or the STFT, the coefficients tend to be highly structured. Nonzero STFT coefficients of audio signals are often clustered near each other (see Figure 5). Similarly, nonzero wavelet coefficients of images are often concentrated close to edges and areas with high-frequency textures. This is apparent in Figure 12.

Thresholding-based denoising can be enhanced by taking into account prior assumptions on coefficient structure. If we have a reason to believe that nonzero coefficients in the signal tend to be close to each other, then we should modify our thresholding scheme. Small coefficients that are isolated probably correspond to noise, and should be suppressed. However, small coefficients in the vicinity of large coefficients may contain useful information and should be left alone. This strategy is known as block thresholding.

Algorithm 4.2 (Denoising via block thresholding). *Let y follow the model in equation (35) and let A be a full-rank linear transformation that sparsifies the signal x . To denoise we:*

1. Compute the coefficients Ay .
2. Apply the hard-block-thresholding operator $\mathcal{B}_\eta : \mathbb{C}^n \rightarrow \mathbb{C}^n$ to the coefficients

$$\mathcal{B}_\eta(v)[j] := \begin{cases} v[j] & \text{if } j \in \mathcal{I}_j \text{ such that } \|v_{\mathcal{I}_j}\|_2 > \eta, \\ 0 & \text{otherwise,} \end{cases} \quad (39)$$

for $1 \leq j \leq N$. The set \mathcal{I}_j is a block of coefficients surrounding index j . The size of each block captures to what extent we expect the coefficients to cluster, and can be adjusted by cross validation. The threshold η can be set according to the standard deviation of Az , or also by cross validation.

3. Compute the estimate by inverting the transform, i.e. setting

$$x_{\text{est}} := L \mathcal{B}_\eta(Ay), \quad (40)$$

where L is a left inverse of A .

Figures 14, 15 and 16 show the results of denoising a speech signal applying block thresholding with blocks of length 5. Figures 17 and 18 show the results of denoising an image by applying block thresholding to its 2D Haar wavelet coefficients, with a block size of 5×5 . As in simple hard thresholding, the threshold is chosen by cross validation on a separate set of signals.

Compared to Wiener filtering, STFT or wavelet thresholding tends to preserve high-frequency features better, as it can adapt to the local structure of the noisy signal. Hard thresholding tends to produce artifacts due to particularly large noise coefficients. In audio, these coefficients can be heard as *musical noise*. Block thresholding alleviates this problem, since it is very rare for neighboring noisy coefficients to be large at the same time.

References

- [1] S. Mallat. *A wavelet tour of signal processing: the sparse way*. Academic press, 2008.

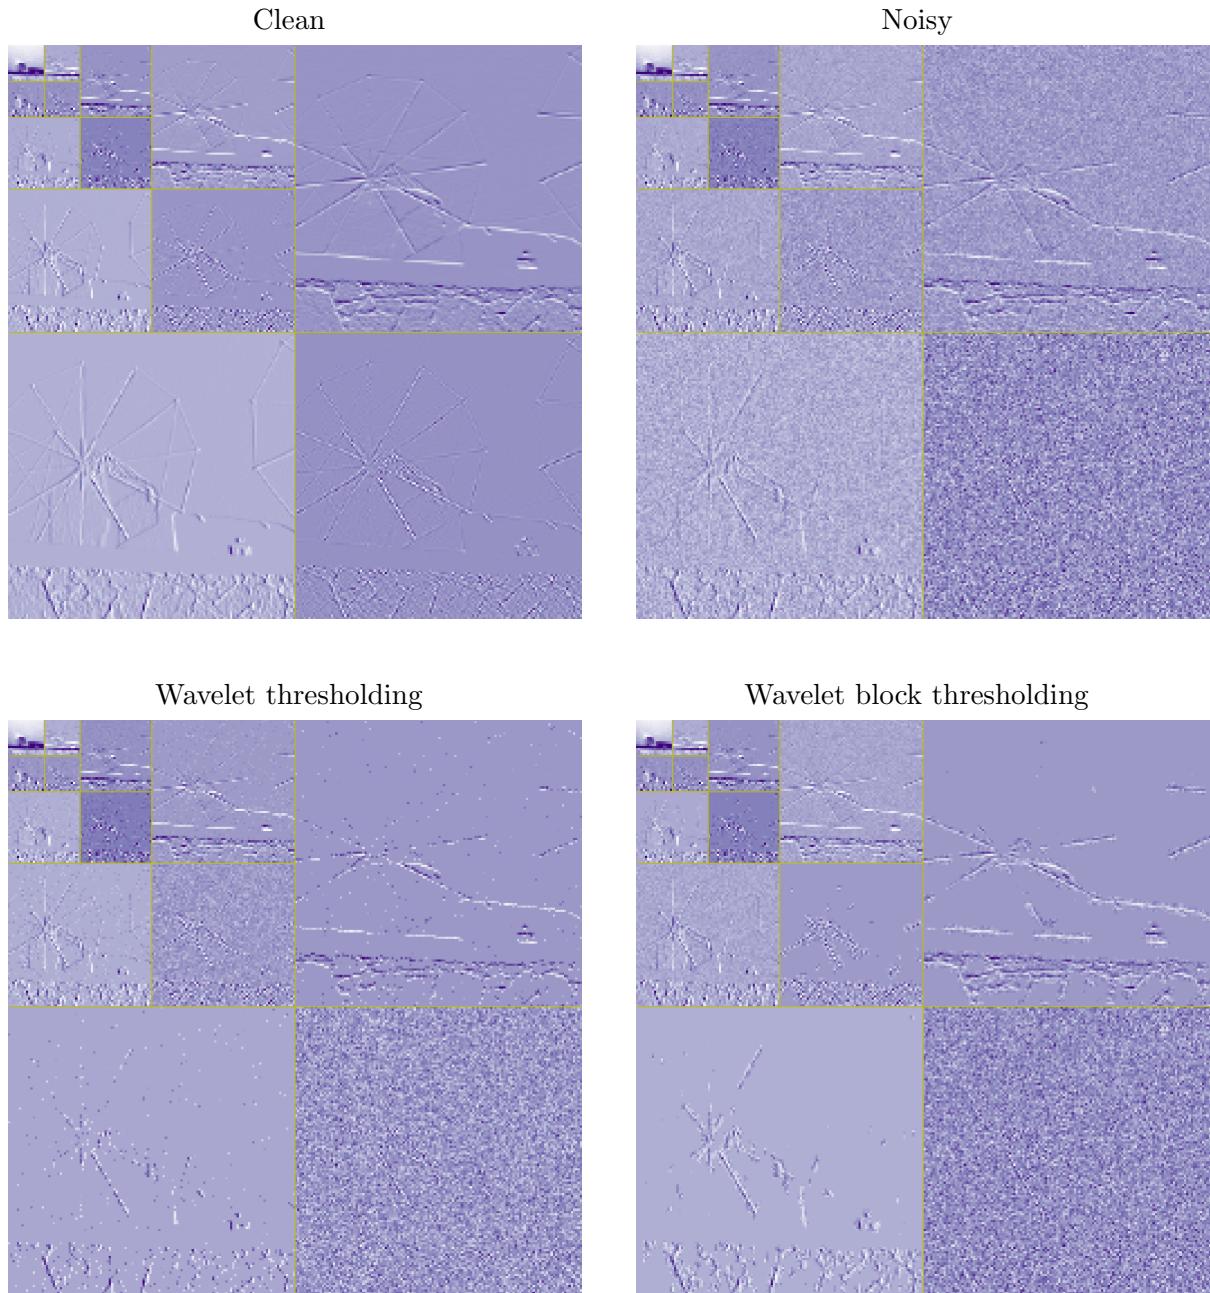


Figure 17: The top row shows the 2D Haar wavelet coefficients of a natural image before and after being corrupted by additive iid Gaussian noise with standard deviation $\sigma = 0.04$. The bottom row shows the result of applying thresholding and block thresholding to the coefficients. The coefficients are grouped by their scale (which is decreasing as we move down and to the right) and arranged in two dimensions, according to the location of the corresponding shifted wavelet with respect to the image.

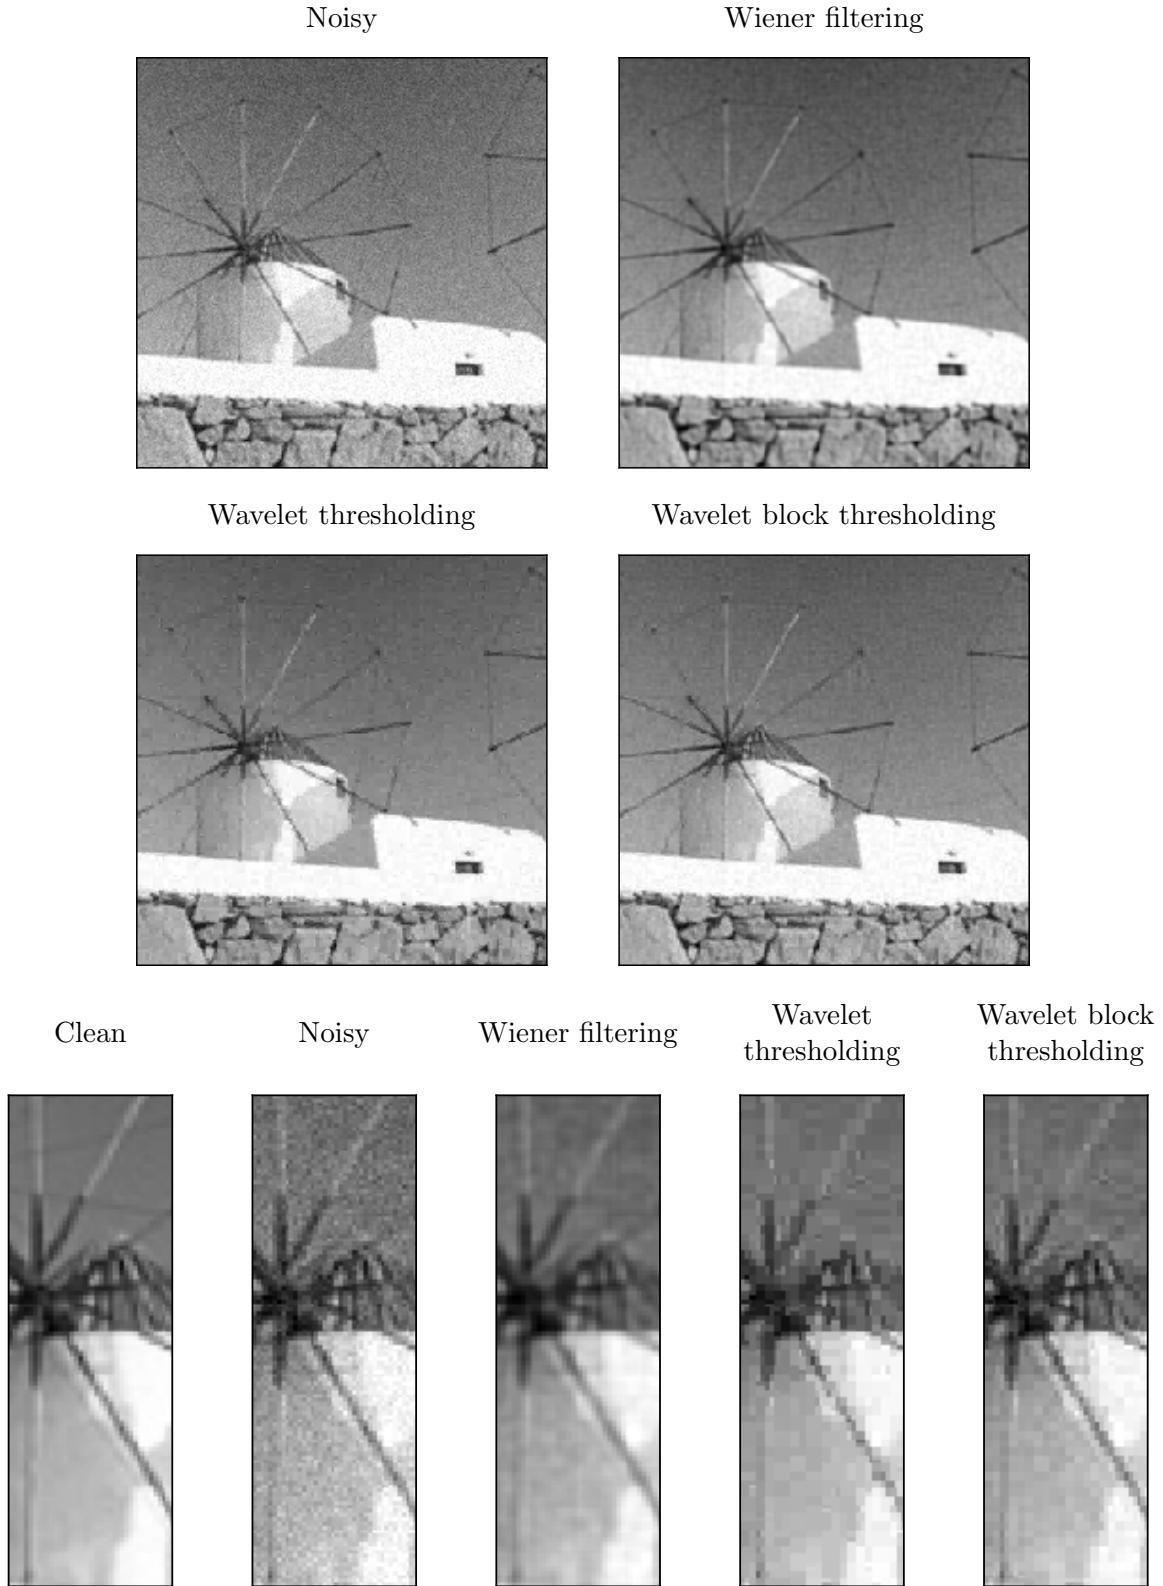


Figure 18: Images corresponding to the coefficients in Figure 17. In addition, the result of applying Wiener filtering to the image is shown for comparison.